

Networks Lab 5 Report

- The topology as you were able to derive it
 - IP addresses of all routers
 - Hosts/ IPs in the Autonomous Systems

```
mininet> nodes
available nodes are:
R1 R2 R3 R4 c0 h11 h12 h13 h21 h22 h23 h31 h32 h33 h41 h42 h43
mininet> net
h11 h11-eth0:R1-eth1
h12 h12-eth0:R1-eth2
h13 h13-eth0:R1-eth3
h21 h21-eth0:R2-eth1
h22 h22-eth0:R2-eth2
h23 h23-eth0:R2-eth3
h31 h31-eth0:R3-eth1
h32 h32-eth0:R3-eth2
h33 h33-eth0:R3-eth3
h41 h41-eth0:R4-eth1
h42 h42-eth0:R4-eth2
h43 h43-eth0:R4-eth3
R1 R1-eth1:h11-eth0 R1-eth2:h12-eth0 R1-eth3:h13-eth0 R1-eth4:R2-eth4 R1-eth5:R4-eth4
R2 R2-eth1:h21-eth0 R2-eth2:h22-eth0 R2-eth3:h23-eth0 R2-eth4:R1-eth4 R2-eth5:R3-eth4
R3 R3-eth1:h31-eth0 R3-eth2:h32-eth0 R3-eth3:h33-eth0 R3-eth4:R2-eth5 R3-eth5:R4-eth5
R4 R4-eth1:h41-eth0 R4-eth2:h42-eth0 R4-eth3:h43-eth0 R4-eth4:R1-eth5 R4-eth5:R3-eth5
c0
```

Figure 1: nodes and net information

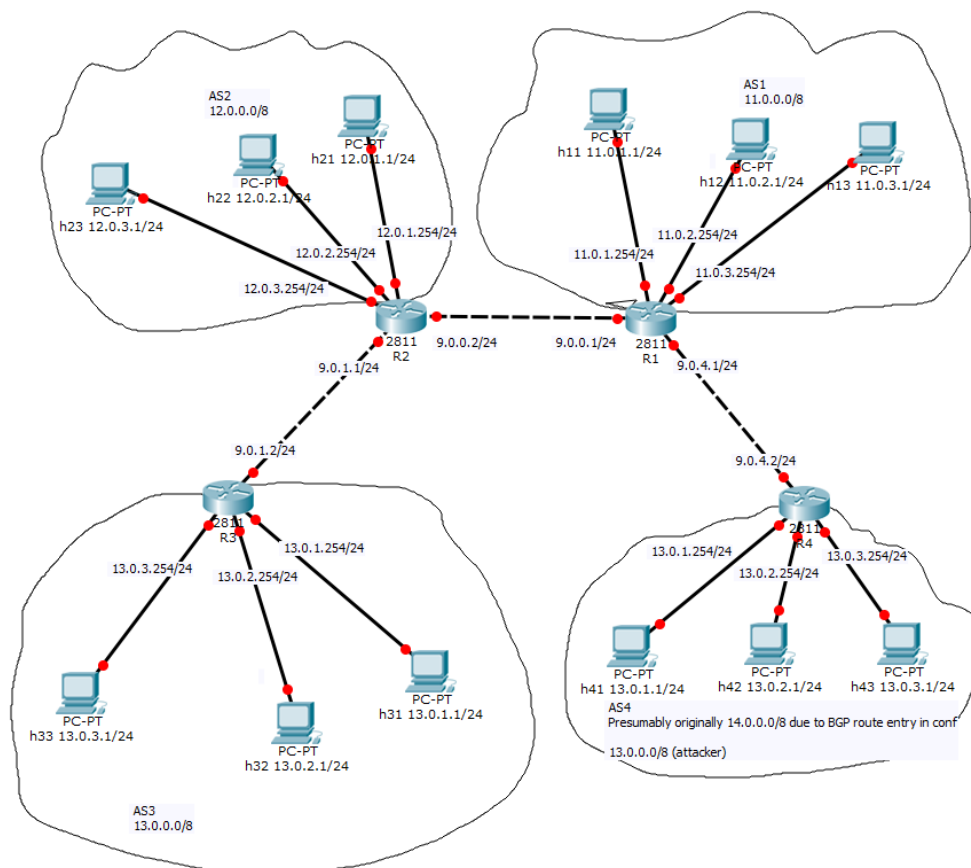


Figure 2: Topology of the Network

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

Using *ifconfig* and *show ip bgp* on each of the routers to derive the following information:

AS1

AS1: 11.0.0.0/8
R1-eth1(h11): 11.0.1.254/24
R1-eth2(h12): 11.0.2.254/24
R1-eth3(h13): 11.0.3.254/24
R1-eth4(R2): 9.0.0.1/24
R1-eth5(R4): 9.0.4.1/24

AS2

AS2: 12.0.0.0/8
R2-eth1(h21): 12.0.1.254/24
R2-eth2(h22): 12.0.2.254/24
R2-eth3(h23): 12.0.3.254/24
R2-eth4(R1): 9.0.0.2/24
R2-eth5(R3): 9.0.1.1/24

AS3

AS3: 13.0.0.0/8
R3-eth1(h31): 13.0.1.254/24
R3-eth2(h32): 13.0.2.254/24
R3-eth3(h33): 13.0.3.254/24
R3-eth4(R2): 9.0.1.2/24

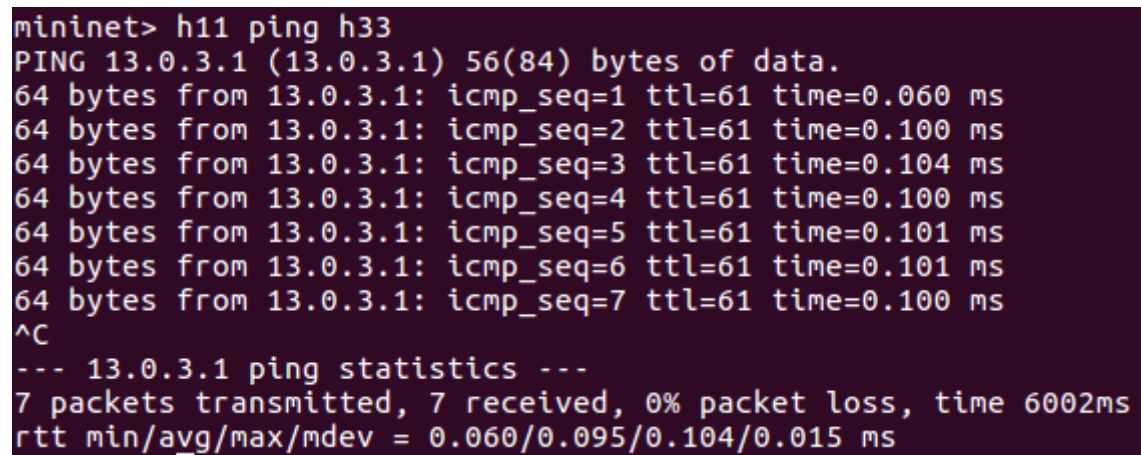
AS4

AS4: 14.0.0.0/8 (to be changed later to spoof as AS3 (13.0.0.0/8))
R4-eth1(h41): 13.0.1.254/24
R4-eth2(h42): 13.0.2.254/24
R4-eth3(h43): 13.0.3.254/24
R4-eth4(R1): 9.0.4.2/24

- **Why was it initially not possible to reach 13.0.1.1 from AS1? How did you find out/what did you do to fix this?**

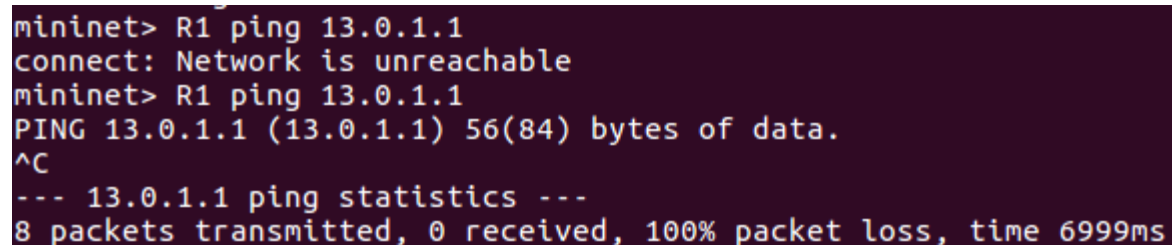
It was initially not possible to reach 13.0.1.1 from R1 AS1 because the packets sent from R1 is able to route to 13.0.1.1 but from R3, it does not know how to route back to R1 AS1 due to a missing routing table entry in R3 which is supposed to tell R3 how to route packets that have destination IP address of R1 (in this case: 9.0.0.0/24 network). There is no routing entry in R3 to route the packets to R1 but there is a routing entry in R3 to route the packets to the host h11 in AS1.

This is verified by using the command *h11 ping h33* (refer to Figure 3), which shows that h11 in AS1 is able to contact h33 in AS3 and that R3 has no problems finding the route to h11 (but not R1). Also, the command *R1 ping 13.0.1.1* (refer to Figure 4) is unable to receive any replies.



```
mininet> h11 ping h33
PING 13.0.3.1 (13.0.3.1) 56(84) bytes of data.
64 bytes from 13.0.3.1: icmp_seq=1 ttl=61 time=0.060 ms
64 bytes from 13.0.3.1: icmp_seq=2 ttl=61 time=0.100 ms
64 bytes from 13.0.3.1: icmp_seq=3 ttl=61 time=0.104 ms
64 bytes from 13.0.3.1: icmp_seq=4 ttl=61 time=0.100 ms
64 bytes from 13.0.3.1: icmp_seq=5 ttl=61 time=0.101 ms
64 bytes from 13.0.3.1: icmp_seq=6 ttl=61 time=0.101 ms
64 bytes from 13.0.3.1: icmp_seq=7 ttl=61 time=0.100 ms
^C
--- 13.0.3.1 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6002ms
rtt min/avg/max/mdev = 0.060/0.095/0.104/0.015 ms
```

Figure 3: h11 ping h33



```
mininet> R1 ping 13.0.1.1
connect: Network is unreachable
mininet> R1 ping 13.0.1.1
PING 13.0.1.1 (13.0.1.1) 56(84) bytes of data.
^C
--- 13.0.1.1 ping statistics ---
8 packets transmitted, 0 received, 100% packet loss, time 6999ms
```

Figure 4: R1 unable to contact 13.0.1.1 at first

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

After configuring R3 to have a routing entry to R1 using *route* (refer to Figure 5), now we are able to successfully ping 13.0.1.1 from R1 and get a response, meaning that the ICMP packets from R1 are now able to recognise the way back and provide a reply.

```
mininet> R3 route add -net 9.0.0.0 netmask 255.255.255.0 gw 9.0.1.1
mininet> R1 ping 13.0.1.1
PING 13.0.1.1 (13.0.1.1) 56(84) bytes of data.
64 bytes from 13.0.1.1: icmp_seq=1 ttl=62 time=0.049 ms
64 bytes from 13.0.1.1: icmp_seq=2 ttl=62 time=0.086 ms
64 bytes from 13.0.1.1: icmp_seq=3 ttl=62 time=0.099 ms
64 bytes from 13.0.1.1: icmp_seq=4 ttl=62 time=0.092 ms
64 bytes from 13.0.1.1: icmp_seq=5 ttl=62 time=0.093 ms
64 bytes from 13.0.1.1: icmp_seq=6 ttl=62 time=0.094 ms
64 bytes from 13.0.1.1: icmp_seq=7 ttl=62 time=0.092 ms
^C
--- 13.0.1.1 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6001ms
rtt min/avg/max/mdev = 0.049/0.086/0.099/0.017 ms
```

Figure 5: After configuring R3, now R1 is able to contact 13.0.1.1

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

- Describe the BGP traffic you were able to observe during re-establishment of routes.

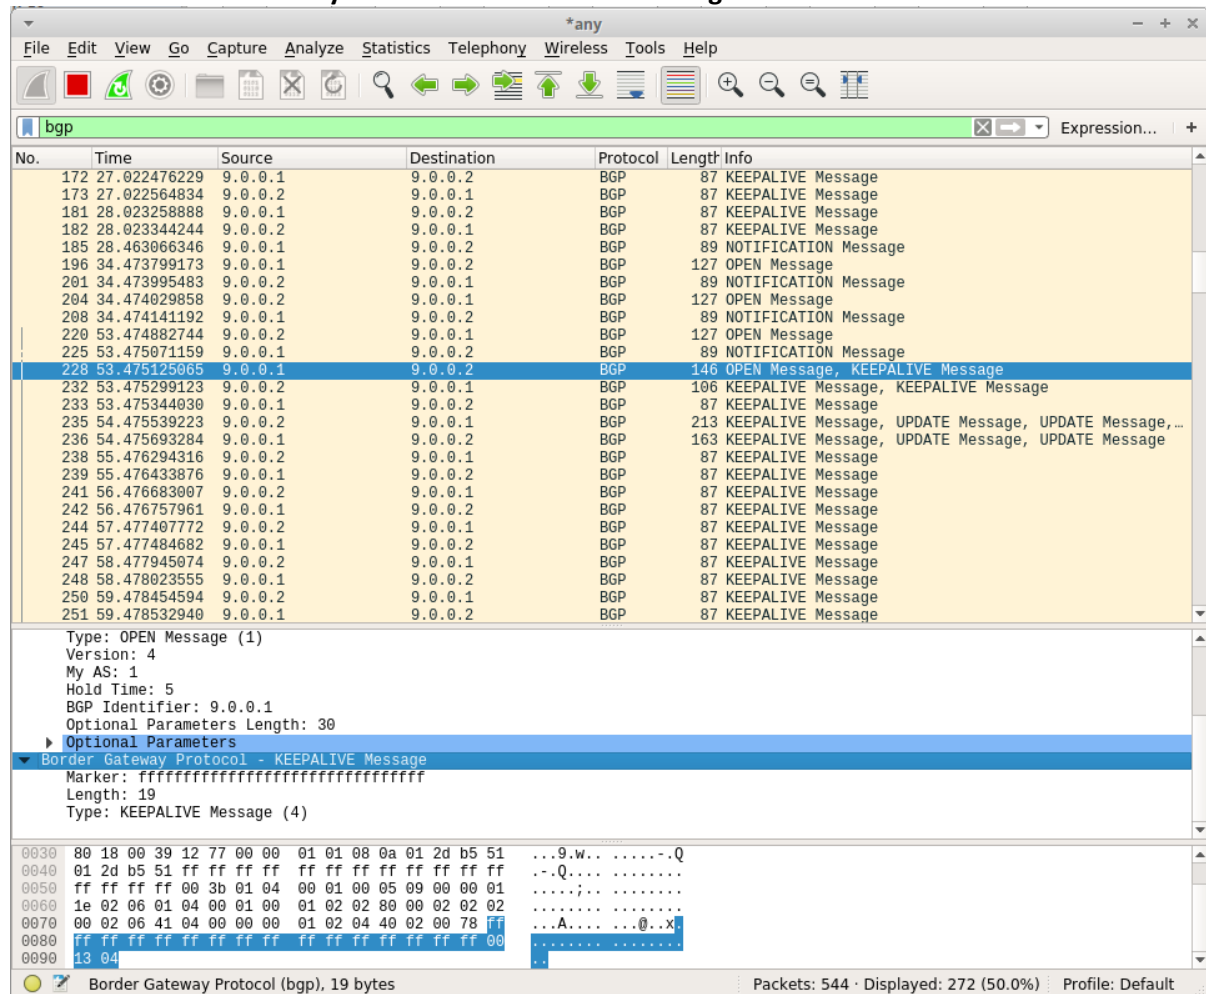


Figure 6: Wireshark packets showing the BGP messages during re-establishment

During re-establishment of routes, it can be seen that there are NOTIFICATION and OPEN messages, followed by subsequent KEEPALIVE messages after the routes have been re-established (refer to Figure 6).

Notification messages is sent by the BGP system when an error condition is detected, and after the message is sent, the BGP session and the TCP connection between the BGP systems is closed. Notification messages consist of the BGP header plus the error code and subcode, and data that describes the error.

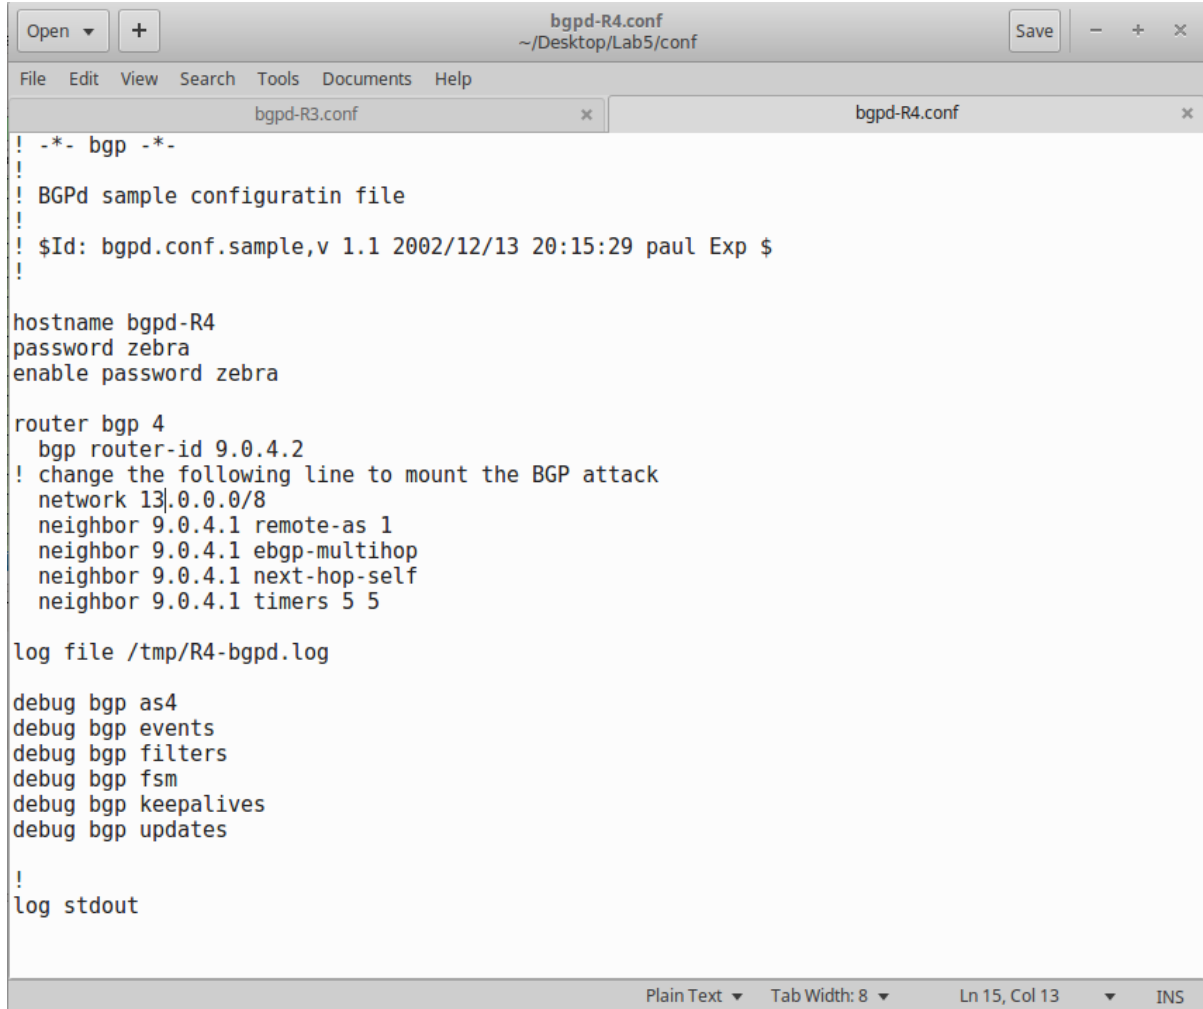
BGP open messages are exchanged after a TCP connection is established between the two BGP systems. Once the connection is established, the two systems can exchange BGP messages and data traffic.

BGP systems exchange keepalive messages to determine whether a link or host has failed or is no longer available. Keepalive messages are exchanged often enough so that the hold timer does not expire. These messages consist only of the BGP header.

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

- **Describe in detail what happened when you started the attack on BGP.**

In order to start the attack, I first modified the bgpd-R4.conf file to change its network parameter from 14.0.0.0/8 to 13.0.0.0/8 (refer to Figure 7 below), which the latter is the webserver we want to disguise as.



```
! *- bgp *-  
!  
! BGPd sample configuratin file  
!  
! $Id: bgpd.conf.sample,v 1.1 2002/12/13 20:15:29 paul Exp $  
!  
  
hostname bgpd-R4  
password zebra  
enable password zebra  
  
router bgp 4  
  bgp router-id 9.0.4.2  
  ! change the following line to mount the BGP attack  
  network 13.0.0.0/8  
  neighbor 9.0.4.1 remote-as 1  
  neighbor 9.0.4.1 ebgp-multihop  
  neighbor 9.0.4.1 next-hop-self  
  neighbor 9.0.4.1 timers 5 5  
  
log file /tmp/R4-bgpd.log  
  
debug bgp as4  
debug bgp events  
debug bgp filters  
debug bgp fsm  
debug bgp keepalives  
debug bgp updates  
  
!  
log stdout
```

Figure 7: Modified bgpd-R4.conf file

```
Wed Oct 12 01:05:03 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:04 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:05 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:06 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:07 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:08 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:09 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:10 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:11 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:12 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:14 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:15 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:16 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:17 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:18 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:19 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:20 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:21 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:22 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:23 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:25 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:26 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:27 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:28 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:29 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:30 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:31 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:32 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:33 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:34 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:36 SGT 2016 -- <h1>*** Attacker web server ***</h1>
Wed Oct 12 01:05:37 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:38 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:39 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:40 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:41 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:42 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:43 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:44 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:46 SGT 2016 -- <h1>Default web server</h1>
Wed Oct 12 01:05:47 SGT 2016 -- <h1>Default web server</h1>
```

Figure 8: website.sh script running

I then ran the provided website script using `./website.sh R1`, which left it continuously contacting a webserver on 13.0.1.1 from R1. It started to return 'Default web server', which is expected since I haven't started the attack. Then, I ran `./start_rogue.sh` to launch the attack. The website results now return '*** Attacker web server ***' (coming from R4) instead of 'Default web server' (coming from R3). Refer to Figure 8 above.

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

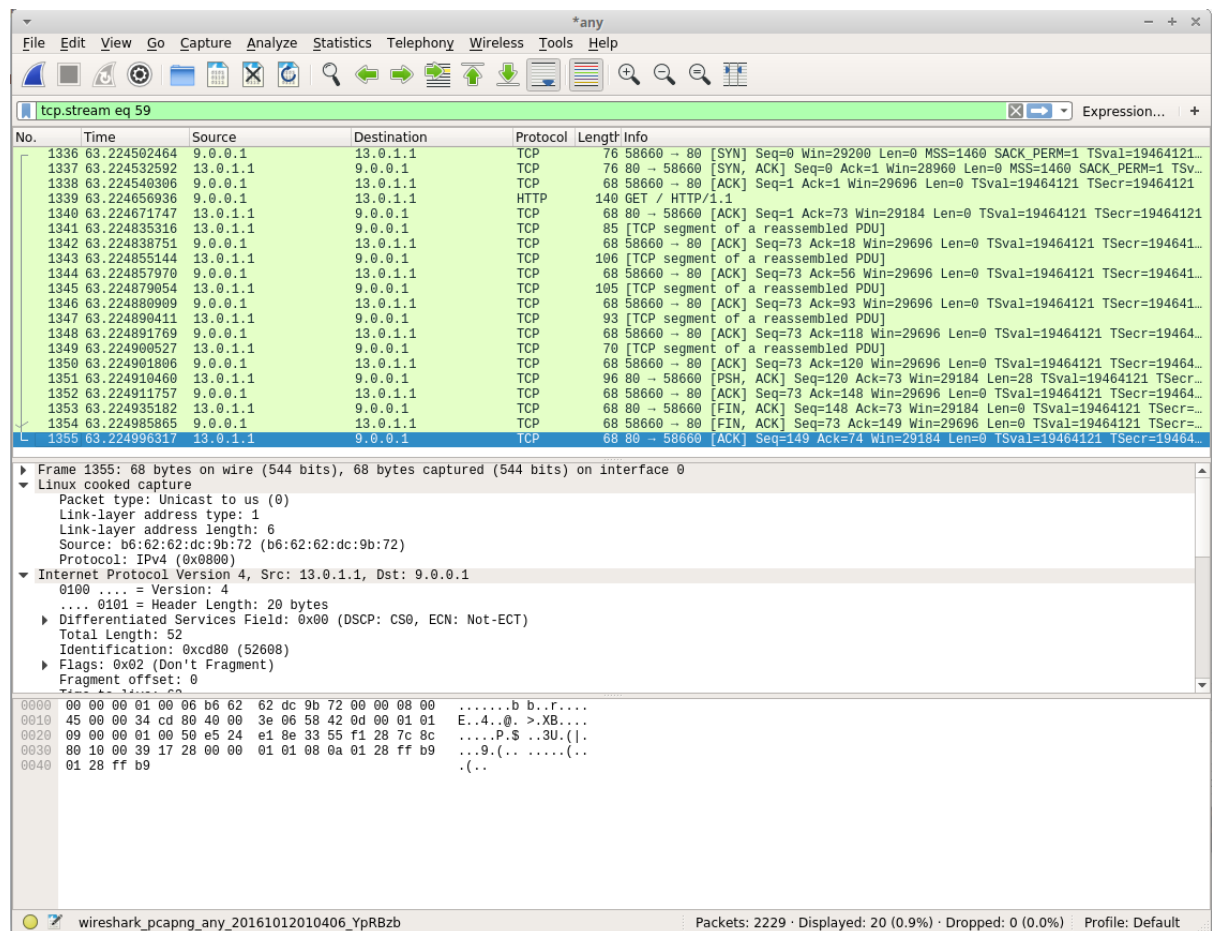


Figure 9: Wireshark packets showing default web server packets

By running Wireshark alongside the attack, I was able to capture and observe the network traffic.

Before I launched the attack, I followed the TCP stream and found out that the network traffic was as desired by the user, with the packets being forwarded from 9.0.0.1 at R1 (refer to Figure 2) to the real destination 13.0.1.1 at AS3. It can also be observed that the source MAC address is b6:62:62:dc:9b:72. Also, I can see from the TCP stream that the message being sent back to R1 was Default web server. Refer to Figure 9 above.

Tan Shun Yu 1001171 ISTD
Nigel Leong 1001095 ISTD

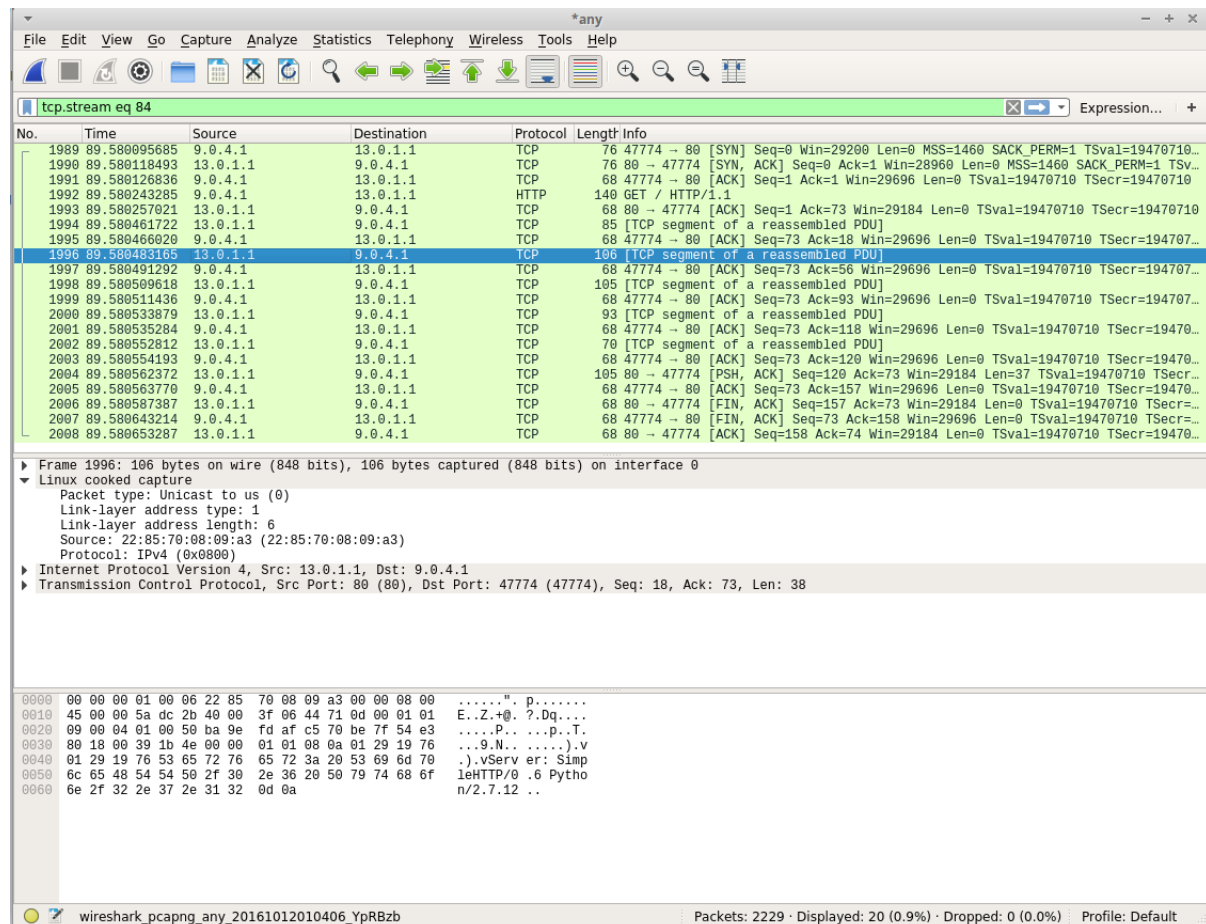


Figure 10: Wireshark packets showing attacker web server packets

After launching the attack, I followed the TCP stream and realised a few differences that confirmed the success of my attack.

The packets are now being forwarded from 9.0.4.1 (refer to Figure 2), which the direction is headed for the attacker AS4 to the spoofed 13.0.1.1 destination. Also, we can observe the source MAC address is now 22:85:70:08:09:a3, which is different from the previous legitimate MAC address of b6:62:62:dc:9b:72. The TCP Stream shows that the message being sent to R1 is now *** Attacker web server***. Refer to Figure 10 above.